

Introduction to Optimization

16th April 2018

Adapted from slides provided by Prof. Jihun Hamm.



Outline

What is optimization?

Convex optimization

- Convex sets

- Convex functions

- Convex optimization

Unconstrained optimization

- Gradient descent

- Newton's method

- Batch vs online learning

- Stochastic Gradient Descent

Constrained optimization

- Lagrange duality

- SVM in primal and dual forms

- Constrained methods



What is optimization?

- Finding (one or more) minimizer of a function subject to constraints

$$\begin{aligned}
 &\underset{x}{\operatorname{argmin}} && f_0(x) \\
 &\text{s.t.} && f_i(x) \leq 0, i = \{1, \dots, k\} \\
 &&& h_j(x) = 0, j = \{1, \dots, l\}
 \end{aligned} \tag{1}$$



What is optimization?

- Finding (one or more) minimizer of a function subject to constraints

$$\begin{aligned}
 &\underset{x}{\operatorname{argmin}} && f_0(x) \\
 &\text{s.t.} && f_i(x) \leq 0, i = \{1, \dots, k\} \\
 &&& h_j(x) = 0, j = \{1, \dots, l\}
 \end{aligned} \tag{1}$$

- Most of the machine learning problems are, in the end, optimization problems.



Examples

- (Soft) Linear SVM

$$\begin{aligned} \underset{w}{\operatorname{argmin}} \quad & \sum_{i=1}^n \|w\|^2 + C \sum_{i=1}^n n \epsilon_i \\ \text{s.t.} \quad & 1 - y_i x_i^T w \leq \epsilon_i \\ & \epsilon_i \geq 0 \end{aligned} \tag{2}$$

- Maximum Likelihood

$$\underset{\theta}{\operatorname{argmax}} \sum_{i=1}^n \log p_{\theta}(x_i) \tag{3}$$

- K-means

$$\underset{\mu_1, \mu_2, \dots, \mu_k}{\operatorname{argmin}} \quad J(\mu) = \sum_{j=1}^k \sum_{i \in C_j} \|x_i - \mu_j\|^2 \tag{4}$$



Outline

What is optimization?

Convex optimization

- Convex sets

- Convex functions

- Convex optimization

Unconstrained optimization

- Gradient descent

- Newton's method

- Batch vs online learning

- Stochastic Gradient Descent

Constrained optimization

- Lagrange duality

- SVM in primal and dual forms

- Constrained methods



Convex sets

Definition

A set $C \subseteq \mathbb{R}^n$ is convex if for $x, y \in C$ and any $\alpha \in [0, 1]$, $\alpha x + (1 - \alpha)y \in C$.

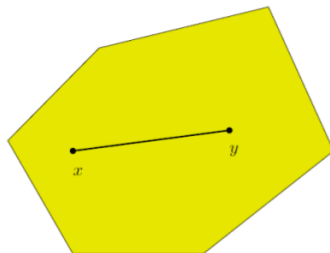


Figure: Convex Set



Convex sets

Example

- All of \mathbb{R}^n
- Non-negative orthant, \mathbb{R}_+^n : let $x \geq 0, y \geq 0$, clearly $\alpha x + (1 - \alpha)y \geq 0$.
- Affine subspaces: $Ax = b, Ay = b$, then

$$A(\alpha x + (1 - \alpha)y) = \alpha Ax + (1 - \alpha)Ay = b$$

.

- Arbitrary intersections of convex sets: let C_i be convex for $i \in \mathcal{I}$, $C = \bigcap_i C_i$, then

$$x \in C, y \in C \Rightarrow \alpha x + (1 - \alpha)y \in C_i \subseteq C, \forall i \in \mathcal{I}$$

.



Convex functions

Definition

A function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is convex if for $x, y \in \text{dom } f$ and any $a \in [0, 1]$,

$$f(ax + (1 - a)y) \leq af(x) + (1 - a)f(y)$$

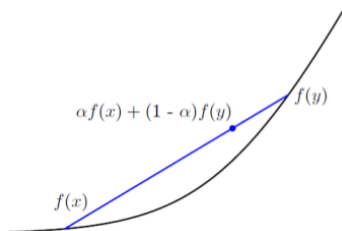


Figure: Convex Function

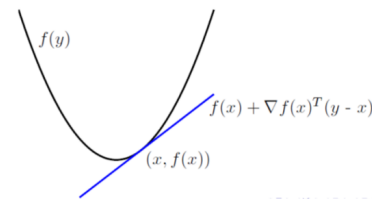


Convexity condition 1

Theorem

Suppose $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable. Then f is convex if and only if for all $x, y \in \text{dom } f$.

$$f(y) \geq f(x) + \nabla f(x)^T (y - x)$$





Subgradient

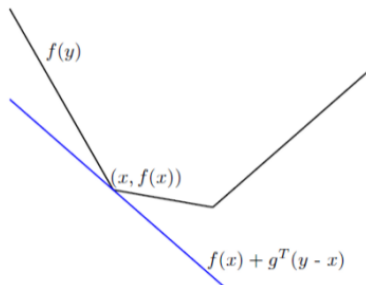
Definition

The subgradient set, or subdifferential set, $\partial f(x)$ of f at x is

$$\partial f(x) = \{g : f(y) \geq f(x) + g^T(y - x) \quad \forall y\}$$

Theorem

$f: \mathbb{R}^n \rightarrow \mathbb{R}$ is convex iff it has non-empty subdifferential set everywhere.






Convexity condition 2

Theorem

Suppose $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is twice differentiable, Then f is convex iff for all $x \in \text{dom } f$,

$$\nabla^2 f(x) \succeq 0.$$





Examples of convex functions

- Linear/affine functions: $f(x) = b^T x + c$
- Quadratic function: $f(x) = \frac{1}{2} x^T A x + b^T x + c$, for $A \succeq 0$.
e.g., for regression:

$$\frac{1}{2} \|\mathbf{X}w - y\|^2 = \frac{1}{2} w^T \mathbf{X}^T \mathbf{X} w - y^T \mathbf{X} w + \frac{1}{2} y^T y$$

.



Examples of convex functions

- Norms (like l_1 or l_2 for regularization):

$$\|ax + (1 - a)y\| \leq \|ax\| + \|(1 - a)y\| = a\|x\| + (1 - a)\|y\|$$

.

- Composition with an affine function $f(Ax + b)$:

$$\begin{aligned} f(A(ax + (1 - a)y) + b) &= f(a(Ax + b) + (1 - a)(Ay + b)) \\ &\leq af(Ax + b) + (1 - a)f(Ay + b) \end{aligned}$$

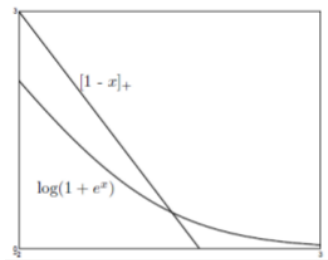
- Log-sum-exp (via $\nabla^2 f(x)$ PSD):

$$f(x) = \log \left(\sum_{i=1}^n \exp(x_i) \right)$$



Examples in machine learning

- SVM loss: $f(w) = [1 - y_i x_i^T w]_+$
- Binary logistic loss: $f(w) = \log(1 + \exp(-y_i x_i^T w))$





Convex optimization

Definition

An optimization problem is convex if its objective is a convex function. The inequality constraints f_j are convex, and the equality constraints h_j are affine.


$$\begin{aligned}
 \min_x \quad & f_0(x) && \text{(Convex function)} \\
 \text{s.t.} \quad & f_i(x) \leq 0, i = \{1, \dots, k\} && \text{(Convex sets)} \\
 & h_j(x) = 0, j = \{1, \dots, l\} && \text{(Affine)}
 \end{aligned} \tag{5}$$



Convex Problems are nice ...

Theorem

If \hat{x} is a local minimizer of a convex optimization problem, it is a global minimizer.





For smooth functions

Theorem

- $\nabla f(x) = 0$. We have $f(y) \geq f(x) + \nabla f(x)^T(y - x) = f(x)$.
- $\nabla f(x) \neq 0$. There is a direction of descent.

Outline

What is optimization?

Convex optimization

- Convex sets

- Convex functions

- Convex optimization

Unconstrained optimization

- Gradient descent

- Newton's method

- Batch vs online learning

- Stochastic Gradient Descent

Constrained optimization

- Lagrange duality

- SVM in primal and dual forms

- Constrained methods

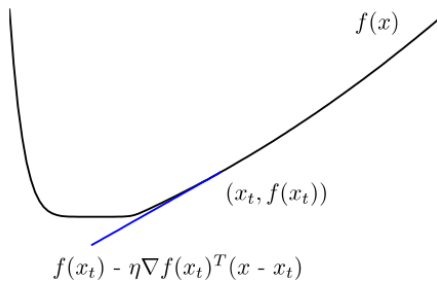


Gradient descent

- Consider convex and unconstrained optimization.
- Solve $\min_x f(x)$.
- One of the simplest approach:
 - For $t = 1, \dots, T$, $x_{t+1} \leftarrow x_t - \eta_t \nabla f(x_t)$
 - Until convergence
 - η_t is called step-size of learning rate.




Single step in gradient descent





Full gradient descent

$$f(x) = \log(\exp(x_1 + 3x_2 - .1) + \exp(x_1 - 3x_2 - .1) + \exp(-x_1 - .1))$$





How to choose step size?

- Idea 1: exact line search

$$\eta_t = \underset{\eta}{\operatorname{argmin}} f(x - \eta \nabla f(x))$$

Too expensive to be practical.

- Idea 2: backtracking (Armijo) line search. Let $\alpha \in (0, 1/2), \beta \in (0, 1)$. Multiply $\eta = \beta \eta$ until

$$f(x - \eta \nabla f(x)) \leq f(x) - \alpha \eta \|\nabla f(x)\|^2$$

Works well in practice.



Newton's method

Idea: use a second-order approximation to function.

$$f(x + \Delta x) \approx f(x) + \nabla f(x)^T \Delta x + 1/2 \Delta x^T \nabla^2 f(x) \Delta x$$

Choose Δx to minimize above:

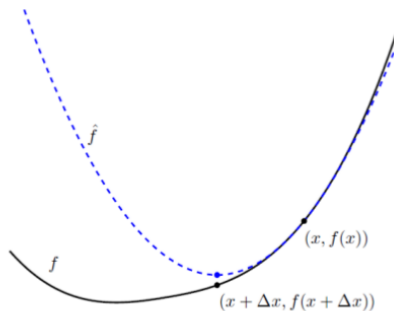
$$\Delta x = -[\nabla^2 f(x)]^{-1} \nabla f(x)$$

This is descent direction:

$$\nabla f(x)^T \Delta x = -\nabla f(x)^T [\nabla^2 f(x)]^{-1} \nabla f(x) \leq 0$$



Single step in Newton's method



\hat{f} is 2nd-order approximation, f is true function.



Convergence rate

- Strongly convex case: $\nabla^2 f(x) \succeq mI$, then "Linear convergence". For some $\gamma \in (0, 1)$, $f(x_t) - f(x^*) \leq \gamma^t$, $\gamma \leq 1$.

$$f(x_t) - f(x^*) \leq \gamma^t, t \geq \frac{1}{\gamma} \log \frac{1}{\epsilon} \Rightarrow f(x_t) - f(x^*) \leq \epsilon$$

.

- Smooth case: $\|\nabla f(x) - \nabla f(y)\| \leq C\|x - y\|$.

$$f(x_t) - f(x^*) \leq \frac{K}{t^2}$$

- Newton's method often is faster, especially when f has "long valleys".



Newton's method

- Inverting a Hessian is very expensive: $O(d^3)$
- Approximate inverse Hessian: BFGS, Limited-memory BFGS
- Or use Conjugate Gradient Descent.
- For unconstrained problems, you can use these off-the-shelf optimization methods
- For unconstrained non-convex problems, these methods will find local optima



Optimization for machine learning

- Goal of machine learning
 - Minimize expected loss $L(h) = \mathbf{E}[\text{loss}(h(x), y)]$ given samples $(x_i, y_i), i = 1, 2, \dots, m$
 - But we don't know $P(x, y)$, nor can we estimate it well.
- Empirical risk minimization
 - Substitute sample mean for expectation.
 - Minimize empirical loss: $L(h) = 1/n \sum_i \text{loss}(h(x_i), y_i)$
 - *a.k.a.* Sample Average Approximation.



Batch gradient descent

Minimize empirical loss, assuming it's convex and unconstrained

- Gradient descent on the empirical loss:
- At each step,

$$w^{k+1} \leftarrow w^k - \eta_t \left(\frac{1}{n} \sum_{i=1}^n \frac{\partial L(w, x_i, y_i)}{\partial w} \right)$$

- Note: at each step, gradient is the average of the gradient for all samples ($i = 1, \dots, n$).
- Very slow when n is very large.



Stochastic Gradient Descent

- Alternative: compute gradient from just one (or a few samples)
- Known as SGD: At each step,

$$w^{k+1} \leftarrow w^k - \eta_t \frac{\partial L(w, x_i, y_i)}{\partial w}$$

(choose one sample i and compute gradient for that sample only)

- the gradient of one random sample is not the gradient of the objective function.
- Q1: Would this work at all?
- Q2: How good is it?



Stochastic Gradient Descent

- Alternative: compute gradient from just one (or a few samples)
- Known as SGD: At each step,

$$w^{k+1} \leftarrow w^k - \eta_t \frac{\partial L(w, x_i, y_i)}{\partial w}$$

(choose one sample i and compute gradient for that sample only)

- the gradient of one random sample is not the gradient of the objective function.
- Q1: Would this work at all?
- Q2: How good is it?
- A1: SGD converges to not only the empirical loss minimum, but also to the expected loss minimum!
- A2: Convergence (to expected loss) is slow:
 $f(w_t) - E[f(w^*)] \leq O(1/t)$ or $O(1/\sqrt{t})$



Practically speaking ...

- If the training set is small, we should use batch learning using quasi-Newton or conjugate gradient descent.
- If the training set is large, we should use SGD.
- If the size of training set is somewhere in between, we use mini-batch SGD.
- Convergence is very sensitive to learning rate, which needs to be determined by trial-and-error (model selection or cross-validation)

Outline

What is optimization?

Convex optimization

- Convex sets

- Convex functions

- Convex optimization

Unconstrained optimization

- Gradient descent

- Newton's method

- Batch vs online learning

- Stochastic Gradient Descent

Constrained optimization

- Lagrange duality

- SVM in primal and dual forms

- Constrained methods



Lagrangian function

Start with optimization Problem:

$$\begin{aligned}
 \min_x \quad & f_0(x) \\
 \text{s.t.} \quad & f_i(x) \leq 0, i = \{1, \dots, k\} \\
 & h_j(x) = 0, j = \{1, \dots, l\}
 \end{aligned} \tag{6}$$

From *Lagrangian* using Lagrange multipliers $\lambda_i \geq 0, \nu_i \in \mathbb{R}$

$$\mathcal{L}(x, \lambda, \nu) = f_0(x) + \sum_{i=1}^k \lambda_i f_i(x) + \sum_{j=1}^l \nu_j h_j(x) \tag{7}$$



Lagrangian function

Original/primal problem:

$$\begin{aligned} \min_x \quad & f_0(x) \\ \text{s.t.} \quad & f_i(x) \leq 0, i = \{1, \dots, k\} \\ & h_j(x) = 0, j = \{1, \dots, l\} \end{aligned}$$

is equivalent to min-max optimization:

$$\min_x \left[\sup_{\lambda \succeq 0, \nu} \mathcal{L}(x, \lambda, \nu) \right]$$

Why?



Lagrangian function

Original/primal problem:

$$\begin{aligned} \min_x \quad & f_0(x) \\ \text{s.t.} \quad & f_i(x) \leq 0, i = \{1, \dots, k\} \\ & h_j(x) = 0, j = \{1, \dots, l\} \end{aligned}$$

is equivalent to min-max optimization:

$$\min_x \left[\sup_{\lambda \succeq 0, \nu} \mathcal{L}(x, \lambda, \nu) \right]$$

Why?

- consider a two-player game, if player 1 chooses x that violates a constraint $f_1(x) > 0$, player 2 chooses $\lambda_1 \rightarrow \infty$ so that $\mathcal{L}(x, \lambda, \nu) = \dots + \lambda_1 f_1(x) + \dots \rightarrow \infty$
- Therefore, player 1 is forced to satisfy constraints.



Dual function and dual problem

- Dual function:

$$g(\lambda, \nu) = \inf_x \mathcal{L}(x, \lambda, \nu)$$

$$= \inf_x \left\{ f_0(x) + \sum_{i=1}^k \lambda_i f_i(x) + \sum_{j=1}^l \nu_j h_j(x) \right\}$$

- Dual problem:

$$\max_{\lambda \succeq 0, \nu} [\inf_x \mathcal{L}(x, \lambda, \nu)]$$

- Primal problem:

$$\min_x [\sup_{\lambda \succeq 0, \nu} \mathcal{L}(x, \lambda, \nu)]$$

- Q: How are primal and dual solutions related?



Weak duality

Dual function lower-bounds the primal optimal value!

Lemma (Weak Duality)

If $\lambda \succeq 0$, then

$$g(\lambda, \nu) \leq f_0(x^*)$$

Proof.

$$\begin{aligned} g(\lambda, \nu) &= \inf_x \mathcal{L}(x, \lambda, \nu) \leq \mathcal{L}(x^*, \lambda, \nu) \\ &= f_0(x^*) + \sum_{i=1}^k \lambda_i f_i(x^*) + \sum_{j=1}^l \nu_j h_j(x^*) \leq f_0(x^*). \end{aligned}$$



Strong duality

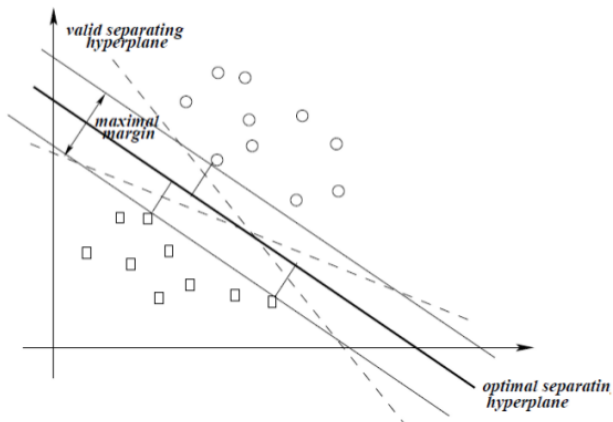
- For convex problems, primal and dual solutions are equivalent!
 $\sup_{\lambda \succeq 0, \nu} g(\lambda, \nu) = f_0(x^*)$
- Equivalently, $\max \min \mathcal{L}(x, \lambda, \nu) = \min \max \mathcal{L}(x, \lambda, \nu)$
- What does the theorem mean in practice?
- When you have a primal constrained minimization problem, which may be hard to solve, you may solve the dual problem, which may be easier to solve (simpler constraints), it yields the same solution!

○○
○○○○○○○
○○○

○○○○
○○○○○
○
○○

○○
○○○●○○
○○

SVM Recap





SVM in primal form

Primal SVM:

$$\begin{aligned} \min \quad & 1/2 \|w\|^2 \\ \text{s.t.} \quad & y_i(wx_i + w_0) \geq 1 \text{ for } i = 1, \dots, m \end{aligned}$$

- for linearly separable cases.
- It is a linearly constrained QP, and therefore a convex problem.



SVM in dual form

The *Lagrangean function* associated to the primal form of the given QP is

$$L_P(w, w_0, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^m \alpha_i (y_i (wx_i + w_0) - 1)$$

with $\alpha_i \geq 0, i = 1, \dots, m$. Finding the minimum of L_P implies

$$\frac{\partial L_P}{\partial w_0} = - \sum_{i=1}^m y_i \alpha_i = 0$$

$$\frac{\partial L_P}{\partial w} = w - \sum_{i=1}^m y_i \alpha_i x_i = 0 \Rightarrow w = \sum_{i=1}^m y_i \alpha_i x_i$$

$$\text{where } \frac{\partial L_P}{\partial w} = \left(\frac{\partial L_P}{\partial w_1}, \dots, \frac{\partial L_P}{\partial w_d} \right)$$

By substituting these constraints into L_P we get its dual form

$$L_D(\alpha) = \sum_{i=1}^m \alpha_i - 1/2 \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j x_i x_j$$



Constrained optimization methods

- Log barrier method
- Projected (sub)gradient
- Interior point method
- Specialized methods
 - SVM: Sequential Minimal Optimization
 - Structured-output SVM: cutting-plane method
- Other optimization not covered in this lecture:
 - Bayesian models: EM, variational methods
 - Discrete optimization
 - Graph optimization



Thank you!